Writing a self-hosting compiler for a purely functional language

(work in progress)

Balázs Kőműves

[bkomuves@gmail.com]

2021.03.24.

What is this?

This is an **experiment** to write a minimalistic self-hosting compiler
in a purely functional language. It's intended both as a learning
experience (for myself) and an experiment in bootstrapping.

Goals:

- ▶ self-hosting compiler
- ▶ very small code-base
- ▶ bootstrappable using a Haskell implementation
- ▶ generated code should have "acceptable" quality

Non-goals:

- ▶ a fully-featured language
- ▶ efficient compilation (the compiler can be slow)
- ▶ bootstrapping "real" compilers from this code base

Why self-host?

Why do bootstrapping / self-hosting?

- ▶ I like simple, self-contained stuff
- ▶ self-hosting is a philosophically desirable property
- ▶ future-proofness: If you can bootstrap from something simple, it won't bit-rot that easily
- ▶ it's a test of the maturity of the compiler
- ▶ the compiler itself is a very good test program!

Exhibit A: Maru, a tiny self-hosting LISP dialect was successfully revived and extended by Attila.

Exhibit B: The York Haskell Compiler (forked from `nhc98`) is now dead, apparently because *nobody knows how to build it anymore*.

The language

The syntax and semantics is based on (strict) Haskell - the idea is
that the compiler should be compilable by both GHC and itself.
This makes development much easier!

The language is basically untyped call-by-value lambda calculus
with data constructors and recursive lets. During parsing, Haskell
features like type annotations, data type declarations, imports etc
are simply ignored - but the compiler itself is a well-typed program!

In fact, a type checker should be probably added, after all; but I
wanted to keep it simple.

Language (non-)features

- no static type system (untyped lambda calculus)
- na data type declarations (constructors are arbitrary capitalized names)
- no module system - instead, we have C-style includes
- strict language (exceptions: if-then-else, logical `and`/`or`)
- ML-style side effects, but wrapped in a monad
- only simple pattern matching + default branch (TODO: nested patterns)
- no infix operators
- list construction syntax `[a,b,c]` is supported
- no indentation syntax (only curly braces), except for top-level blocks
- only line comments, starting at the beginning of the line
- no escaping in character / string constants yet
- (universal polymorphic equality comparison primop)

Compilation pipeline

1. lexing and parsing
2. collecting string literals
3. partition recursive lets into strongly connected components
4. TODO: eliminate pattern matching into simple branching on constructors
5. recognize primops
6. collect data constructors
7. scope checking & conversion to core language
8. inline small functions + some basic optimizations
9. closure conversion
10. TODO: compile to a low-level intermediate language
11. final code generation

The runtime system

The runtime system is relatively straightforward:

▶ there is a heap, which can only contain closures and data constructors

▶ there is stack (separate from the C stack) which contains pointers to the heap

▶ heap pointers are tagged using the lowest 3 bits

▶ heap objects fitting into 61 bits (ints, chars, nullary constructors, static functions) are not allocated, just stored in registers / on stack

▶ statically unknown application is implemented as a runtime primitive

▶ there is a simple copying GC - the GC roots are simply the content of the stack (plus statically known data)

▶ there are some debug features like printing heap objects

▶ primops: integer arithmetic; integer comparison; bitwise operations; lazy `ifte/and/or`; basic IO.

Current state of the experiment

Current state:

- ▶ it can compile itself successfully;
- ▶ the source code is not as nice as it could be;
- ▶ the generated code is pretty horrible;
- ▶ there are some desirable features missing.

Code size:

- ▶ the compiler: ~ 1900 lines (required code)
- ▶ type annotations: $+525$ lines (ignored)
- ▶ the C runtime: ~ 600 lines (including some debug features)

Mistakes I made

I made a huge amount of mistakes during the development, which made it much harder and longer than I expected.

- ▶ Trying to self-host before the compiler actually works. Writing a compiler is hard work, you don't want to do it in a limited language!

- ▶ Not figuring out the precise semantics upfront:
    - ▶ exact details of lazyness vs. strictness
    - ▶ controlling exactly where heap / stack allocations can happen
    - ▶ recursive lets

- ▶ Not having debugging support from early on (for example: good error messages, threading names & source locations, RTS features, pretty printing AST-s)

- ▶ Trying to target assembly first, instead of a simpler target like C; it's just extra cognitive load

- ▶ Generating code directly, without having a low-level IL

- ▶ Using de Bruijn indices (as opposed to levels) in the generated code

Bug gallery

Some examples of bugs which were really painful to debug:

- ▶ IO ordering bugs (GHC *really* likes to optimize my IO implementations away...)

- ▶ GC liveness bugs (see the next slide)

- ▶ stack management, and using de Bruijn indices in the generated code

- ▶ C stack overflow gives you a segfault with no information

- ▶ thunks (there shouldn't be any, this is supposed to be a strict language after all...)

- ▶ recursive lets are a PITA if you allow non-lambdas

- ▶ there was a very simple bug (wrong variable name) in the implementation of generic application in the RTS which was not triggered by compiling the whole compiler, only after I did inlining

- ▶ inlining code which does computation makes the runtime blow up exponentially, which looks *very* similar to an infinite loop...

Example GC bug - Marshalling from C strings

```
        heap_obj marshal_from_cstring(char *cstring) { ...

v1:     obj    = heap_allocate(Cons,2);
        obj[1] = cstring[0];
        obj[2] = marshal_from_cstring(cstring+1);
        return obj;

v2:     rest   = marshal_from_cstring(cstring+1);
        obj    = heap_allocate(Cons,2);
        obj[1] = cstring[0];
        obj[2] = rest;
        return obj;

v3:     rest   = marshal_from_cstring(cstring+1);
                 push(rest);
        obj    = heap_allocate(Cons,2);
        obj[1] = cstring[0];
        obj[2] = pop();
        return obj;
```

Future work

I'm not convinced that it is worth hacking on this toy compiler (as opposed to work on a full-featured "real" compiler), but in any case, here is a TODO list:

- ▶ add nested patterns, and patterns in lambda arguments
- ▶ do some optimizations, so that the generated code is better
- ▶ tail call elimination
- ▶ escaping in string constants
- ▶ compile to a low-level intermediate language first, and do codegen from that
- ▶ add a type system
- ▶ more backends: x86-64 assembly and/or machine code; some simple VM